



Joe Eichenbaum

Partner, 17A

Zack Wortman

Partner, 17A

September 2026

The Only Way to Start Is by Starting

And that's easier than you
think: the first steps of AI
adoption for state and
local government agencies



Executive summary

The pressure to “do something about AI” has arrived, and the instinct is to meet it with a plan, whether a task force, a roadmap, or a year-long strategy. That instinct is one of the main things slowing agencies down. The agencies getting value from AI did not get there by planning. They started small and learned as they went. This paper describes how to do the same.

Three ideas run through it.

- *You are ready.* Whatever shape your agency is in, with old systems, a thin technical bench, and no AI policy, you can get useful value from AI this quarter. The foundation is four things you put in place once, and none of them needs meaningful new budget or new authority.
- *Start by starting.* Don’t spend six months writing a strategy you can’t yet write well. Pick one or two small, low-risk use cases, run them against real work, and let the strategy grow out of what you learn.
- *One size fits one.* The most useful early tools are small and custom, built to solve one team’s specific problem. Examples include synthesizing public comments on a rule, flagging incomplete incident reports as they are filed, and reconciling records across two systems that don’t talk to each other. None of these is an enterprise platform, and AI has made tools like them newly cheap to build.

Finding out costs one accountable owner, which is not a full-time job, five to ten software licenses at roughly \$100 a seat per month, and ninety days. At the end of that time you can have a working tool, a team that trusts it, and a defensible number for what it saved. That is a far better footing for the big AI decisions than another planning cycle.

Why now, when so many other things are urgent?

If you run a state or local government agency, or a division or program inside one, you are already getting the message that you need to act on AI. Vendors are pitching it, your staff are using ChatGPT to draft emails and summarize documents whether or not anyone formally approved it, and a legislator has probably asked what you are doing about it. All of it lands on a desk that has no free time on it. The challenge is to start in a way that gives you time back instead of taking more.

Government agencies are among the hardest places to adopt new technology. Whatever the mission, whether health coverage, public safety, education, or transportation, most agencies run dozens of programs on a patchwork of decades-old case, records, and financial systems, under funding rules that shape how technology gets built and paid for, and with some of the most sensitive data government holds. In most agencies, AI is also nobody’s job. It sits between IT, program, and business leaders, relevant to all of them and owned by none of them.

The most common reason agencies delay on AI is bandwidth. Few agency leaders are skeptical of the tools, but their teams are already running flat out on new federal requirements, legislative mandates, and dozens of other time-sensitive bodies of work. AI does not have to be one more initiative competing for those hours. Used thoughtfully, it is a tool that can give some of that time and attention back in concrete, measurable ways.

There is also a longer-term reality. The capabilities of these tools are changing in ways that will, over time, reshape how a great deal of office work gets done. You do not need a position on that today, but the agencies building hands-on experience now through small, low-risk projects will be far better prepared to make good decisions on the bigger questions as they arrive. Making that progress takes one or two small projects, learning as you go, and scaling what works. A large budget, a dedicated team, and an ambitious roadmap can all wait.

Many government workers aren't behind on AI, but most agencies are

It would be easy to assume government is hopelessly behind the private sector on AI. There are tech-forward corners of the private sector that are *far* ahead, but at the level of the individual worker there is little difference. As of late 2025, public-sector employees were using AI at work at about the same rate as their private-sector peers.¹ Walk the halls of most agencies and you will find staff already leaning on AI for everyday drafting and summarizing, whether or not policy and approved tools have caught up.




Figure 1. Public-sector workers use AI at nearly the same rate as private-sector workers (43% versus 41%). The gap is organizational. Only 37% of public-sector workers say their organization has a clear AI strategy, compared with 53% in the private sector. Source: Gallup, Q4 2025.

The gap is organizational, and it matters. Where leaders set a direction and actively support AI use, employees use it more often and in more ways, and they get more entrepreneurial about finding and sharing ways it fits their own work.

The one exception is also the most useful one. Government's technical staff, the people who run IT, are far behind their private-sector counterparts, and that is an opportunity. The first place AI broke through was coding, and over the past two years it has reshaped how private-sector technology teams work, while most government IT shops have not felt any of it yet. Coding and data work is the most proven thing AI does at scale, which raises a fair question about whether your first pilot should go to a program office at all. It may belong in your own IT team.

When AI use stays ad hoc and individual, the gains stay small and uneven while the risks grow: inconsistent quality, sensitive data flowing through unapproved tools, and lessons that never get shared, with little visibility for the agency to provide oversight. Since many of your staff already use AI, the work is to give that energy direction, guardrails, and a way to build on itself.

Laying the foundation

In practice, *you are ready* means four things, put in place once, none of which requires new authority or meaningful budget. Letting a few curious people experiment in a corner is not enough. These four things put your agency in a position to learn deliberately and quickly, by running small experiments on real work, seeing what works, and building from evidence rather than speculation. That is the difference between scattered individual use and broad, systemic adoption.

- *A shared AI tool you can trust with your work.* Use the business or enterprise tier of a major assistant, Claude (Anthropic) or ChatGPT (OpenAI), and skip the free public version. The paid tiers give your agency admin controls and a contractual commitment that your information will not be used to train the underlying model. We have left Gemini (Google) off this shortlist on purpose. In our testing it does the core work less well, lagging Claude and ChatGPT on the writing, summarizing, and coding these pilots run on, and it is clunkier to use. Its one advantage, tight Google Workspace integration, would not help a Microsoft-shop government anyway. We have worked with all three and currently prefer Claude, but this is a fast-moving field. Choose a tool you can start with. You are not choosing a permanent platform.
- *Enough people with access to test and learn, but no more.* You do not need an agency-wide deployment to start, and you should not wait for one. A team can begin with roughly five to ten licenses. Budget about \$100 per person per month rather than the \$20–30 chat-only tiers. A pilot where people only ever open a chat box will not change how the work gets done. You need staff in the *agentic* tools that can do work alongside them (Claude Code, or at minimum a “Cowork”-style agent), because that is where a manual workflow gets rebuilt. Even so, five to ten seats runs roughly \$6,000 to \$12,000 a year, still under the small-purchase threshold that avoids competitive procurement in many states. If you have an IT governance board, put it on the next agenda. The tool does not need to integrate with your existing systems to deliver value, though integration can come later and open up more use cases.

- *Guardrails that let people use AI on real work safely.* People need to experiment without taking on risk for the agency. On a business or enterprise tier, your inputs stay within a protected environment and are not used to train the model, but the specifics vary by vendor and tier. Look for recognized security authorizations such as FedRAMP or StateRAMP, have your IT and security teams verify that the setup meets your state's data-handling requirements before any real data goes in, and make sure a human verifies AI output before it is used. If your security team's reflex is to say no on principle, the section on barriers below takes that on directly.
- *An owner and a simple, repeatable rhythm of learning.* Name one accountable owner to move this forward. They do not need to build everything or become the agency's AI expert. Their role is to run a simple, repeating loop, shown in Figure 2.




Figure 2. The owner's loop. None of this is a full-time job, especially at the start.

The owner should know going in that the best ideas about what to build rarely come from senior leadership. They come from the front-line staffer or analyst who can see exactly where the work snags. Most of the owner's job is finding those people and giving their ideas somewhere to go.

Choosing and running your first pilots

Putting *start by starting* into practice takes three things: a way to pick well (the checklist below), a sense of how the first few months unfold (the ninety-day walk-through), and a concrete picture of what a first project looks like (the worked examples).

A checklist for picking your first use cases

For your first one or two use cases, choose work that meets most of these criteria. The last two may look contradictory, but they are alternative green lights. Highly manual work shows AI's value most clearly, and work already done in code shows where AI can do the same job faster with less specialized skill.

- *Internal users first.* Serve your own staff before you serve citizens directly. This minimizes risk, and it makes your first users staff who can tell you what is useful, what is not, and how to refine it.

- *Small user base.* It is much easier to build a tool that fits five or ten people well than one for hundreds of front-line staff or dozens of different roles across the agency.
- *Decision support, with a human deciding.* AI helps people do their jobs better and does not replace their judgment. A human stays in the loop.
- *Not highly sensitive.* Avoid legally protected categories (health records, case-level client data, criminal justice information) for your first projects.
- *Currently highly manual.* The more time staff spend today summarizing, reformatting, or searching, the more obvious the value will be.
- *Already done with code or near-code.* If a team currently leans on Python scripts, SQL queries, or complex Excel macros, that is a strong signal that AI can do the same work faster and with less specialized skill.

What the first ninety days can look like

Every agency's path will differ. The following is one illustrative shape, and its value is in showing how small the individual steps are.



Figure 3. The first ninety days at a glance.

Weeks 1–2: Set up. Name the owner, choose the tool, and purchase five to ten licenses. Put the effort on the agenda of your IT governance board or equivalent, and have IT and security confirm the data-handling guardrails. Send the pilot group a one-page memo covering what the tool is for, what data stays out of it for now, and the expectation that a human reviews every output.

Weeks 3–4: Pick the first use case. Ask the pilot group and a few trusted managers a single question: “What recurring work eats the most time for the least judgment?” Collect the answers into a running list and score the candidates against the checklist above. Pick one, name the two or three staff who own that work, and define in one sentence what done looks like, for example “the next public-comment summary is produced in two days instead of two weeks, at a quality the policy director signs off on.” Don’t run this stretch entirely on paper. Spend part of these two weeks prototyping. Put the candidate problems in front of the AI, ask what is possible, what it would take, and what a first version might look like, and then have it sketch a demo or a plan. You will learn more from two days of hands-on work than from two weeks of meetings. Some teams run this stretch as a small internal hackathon, which is the right instinct.

Weeks 5–8: Run it against real work. The team uses the tool on the actual task from the start, since a demo version teaches you very little. The owner runs a thirty-minute check-in each week to cover what worked, what did not, and what would make it better. Expect the first attempts to be mediocre and the fourth or fifth to be useful. That improvement curve is the learning you are buying.

Weeks 9–10: Decide. Hold a retrospective and answer three questions. Did it save time or improve quality? Would the team keep using it tomorrow if nobody made them? What would it take to make it twice as useful? Then decide whether to keep, refine, or stop, and write the answer down in a one-page summary whichever way it goes.

Weeks 11–13: Keep going. Start the second use case from the running list, applying what the first one taught you. Share the one-pager with your leadership team and the governance board. If the first pilot worked, let the team that ran it tell the story, because peer credibility recruits the next round of volunteers better than any directive.

Three illustrative examples

Each example below is *one size fits one* in action, a small, fit-for-purpose tool for one team's workflow.

Example 1: Structured, repeatable synthesis of public comment (difficulty 1 of 5, very easy)

A proposed rule can draw hundreds of comments, and staff must group and respond to the substantive ones before it is finalized. Many staff already use AI to help, by pasting comments into a chat and asking for a summary. A way to make this both more useful and more governable is to give your AI assistant a markdown file of standing instructions that it applies to *any* public comment summary: group comments by theme, count comments by topic, flag comments from specific stakeholders, and put the final deliverable in your standard format. Even when a specific summary needs something different, it is much easier to adjust from a baseline than to explain everything from scratch each time.

- Write a markdown file with standard rules for all summaries.
- Upload the proposed rule change, the public comments, and the markdown file to your AI assistant, which produces a draft summary.
- A staff member reviews the draft, refines it as needed, and approves it.
- Save the refined instruction file, so the next rulemaking starts from a sharper default.

Example 2: A refreshable check on incident report completeness (difficulty 2 of 5, relatively easy)

Much of what happens downstream of a law enforcement incident report depends on the report being complete, including charging decisions, victim follow-up, crime statistics, and state and federal reporting. Records management systems accept whatever gets filed, so gaps such as a blank victim field or a missing offense code

surface weeks later as a rejected state submission, a prosecutor's callback, or a statistic nobody quite trusts. Today that means a supervisor checking reports by hand, if anyone checks at all. A small, reusable tool can do it instead. You describe the job in plain language, and the assistant writes the code.

A note on context: this is where you move from a chat window into an agentic tool like Claude Code. The first thirty to sixty minutes are steep, and worth doing alongside someone who has done it before, but it gets much easier after that, even for people who have never coded. Prompting a coding assistant is also nothing like writing code. There is no syntax that must be perfect. You can describe the problem in your own words, including what you do today, the output you want, and what is painful about the current process.

- Describe the goal, and provide an extract of last month's reports and the fields that must be complete for each incident type.
- The assistant builds a tool that scans every report, flags blank or inconsistent fields, and summarizes the patterns by incident type, unit, and shift.
- A records manager runs it on current data, checks the flags against reports already known to be problems, and refines it until the output is trustworthy.
- The team reruns it nightly or weekly, turning report quality from an end-of-month surprise into feedback a supervisor can act on while the officer still remembers the call.

There is often more than one workable design. The tool could sign into the records system itself, pull data through an API, or take a file someone exports and hand back a finished one. All three are legitimate, so start with whichever is easiest.

Example 3: Reconciling eligibility and enrollment files across systems (difficulty 4 of 5, harder)

One technical person with a little vision can pull this off. Every month, a state Medicaid agency and its managed care plans trade enrollment files, comparing the state's record of who is covered against each plan's. Most records agree, but a few hundred each cycle do not, such as a retroactive disenrollment a plan has not processed, a newborn on a plan's rolls but not yet in the state system, or a member showing active in two plans at once. An analyst works that list in a spreadsheet when there is time, which means the state can keep paying a plan its monthly fee for a member who already left, sometimes for months, and getting that money back afterward is its own small project. The file formats here are Medicaid's, but nearly every agency has a version of the same problem, two systems that are supposed to agree and do not quite.

A small team can build a tool for this in weeks. Using it each month would look like this.

- The analyst hands the tool this month's files from both sides and runs it.
- A few minutes later she has a usable list of every mismatch, sorted by what it will cost if ignored, with each one explained in a sentence a person can act on.

- She resolves them in the source systems, and the ones she cannot resolve she escalates with the explanation already written.
- The next cycle follows the same routine. When a new kind of mismatch shows up, she describes it to the tool in a sentence, and it gets caught from then on.

The work has not disappeared, and the analyst still decides every case. The gain is that the whole list gets worked every cycle instead of whatever fraction someone had time for, and the mistakes get caught in the month they happen rather than a year later in an audit.

How to know if it's working

Measurement here should take minutes and should not require a dashboard. A few simple signals, checked at the week-nine retrospective and periodically afterward, are enough.

- *Hours returned.* Ask the team to estimate how long the task took before and how long it takes now. A rough estimate is fine as long as the direction is right.
- *Cycle time.* How long does it take from when work arrives to when it is done? That is the two-weeks-to-two-days measure, and it is often the number leadership and legislators care about most.
- *Dollars and service quality, where you can name them.* Some wins land straight on the ledger, such as contractor hours you did not buy, overtime you did not pay, and rework you avoided. Others show up as service quality you can count, such as faster turnaround on constituent replies, fewer notices that generate confused phone calls, and lower error or appeal rates. When a pilot moves one of these, it is usually the number that travels furthest with budget offices and legislators.
- *Voluntary repeat use.* Of all these signals, this one is the most reliable. If the team keeps using the tool when nobody is checking, it is working. If usage decays after week three, it is not, no matter what the demo looked like.
- *Quality on spot-check.* Have the relevant supervisor periodically compare AI-assisted output against the old way. The standard is output that is at least as good as the old way and takes less effort to produce.

Give yourself explicit permission to stop. A use case that is not paying off after a fair trial should be shut down, and shutting it down is a success of the process rather than a failure of the program, because you learned something for a few thousand dollars that enterprise programs routinely spend millions to learn. Write down why it did not work and move to the next item on the list. Agencies that never shut anything down end up maintaining a collection of half-used tools, while agencies that stop quickly learn faster.

Failure modes to watch for

Early AI efforts tend to stall from the inside in a handful of recognizable ways, all of them avoidable and most of them visible within a few weeks. External frictions such as procurement, legal, and workforce concerns are covered in the next section.

Pilot purgatory. The demo impresses, victory is declared, and the tool never becomes part of how the work gets done. You can recognize it when nobody can name the staff member who used the tool on real work this week. The ninety-day structure above guards against it by defining done as a real task completed faster, which a demo cannot satisfy.

The tool nobody opens after week three. Initial enthusiasm fades because the tool is almost but not quite fit for the real workflow, and an almost-fit tool loses to habit every time. Decaying voluntary use is the warning sign, and the weekly check-in is where you catch it. Treat “I stopped using it because...” as the most valuable sentence in the pilot, and either refine the tool until the friction is gone or stop.

The single enthusiast. One energetic person builds something useful, then changes roles, and the tool dies with their departure. From the start, require that at least two people can run any tool the team decides to keep and that the instructions live in a shared file where anyone can find them.

Sensitive data creep. The pilot starts with safely scoped data and works well, and staff reasonably start reaching for the more sensitive use cases before the guardrails have caught up. This is a sign of success and also a moment of risk. Make the boundary explicit in the one-page memo, repeat it at check-ins, and route every request to expand scope through privacy and security review. The pilot team should not decide those requests on its own.

Premature enterprise ambition. Six weeks in, momentum builds and someone proposes scaling the comment-synthesis tool to every division, with requirements gathering and a steering committee. Resist the pull. *One size fits one* is the operating principle because fit for purpose is what made the pilot work. Scaling is sometimes the right call, but it should be earned by sustained voluntary use rather than projected from a good first month.

Waiting for the perfect policy. The opposite failure is that nothing starts because the comprehensive AI policy is not finished. A one-page memo naming the approved tool, the data that stays out, and the human review requirement is enough governance for a five-person pilot on non-sensitive work. Let what the pilots teach you inform the full policy.

The veto that never resolves. The tool never gets bought. Security flags a concern, procurement is not sure how to classify the purchase, or legal wants one more review, and because pushing it through is no one’s job, “not yet” gradually becomes “no.” Plenty of efforts die at this stage, before a single pilot runs. Treat the purchase itself as a use case with a named owner and a deadline, and walk into the security conversation prepared. The next section covers how.

Getting past the external barriers

The previous section covered failures that come from inside the effort. This one covers the external frictions you are likely to meet once you are moving. Each can be worked through, and most are easier at pilot scale than your instincts may suggest.

Workforce concerns. The question underneath is almost always “Is this how they cut our jobs?” It deserves a straight answer. These pilots are decision support with a human in the loop, aimed at the parts of jobs people like least, such as summarizing, reformatting, and searching, in agencies that almost always have more work than capacity. Let front-line staff choose and shape the use cases, and let the pilot teams report the results themselves. Staff who got hours back are more persuasive to their colleagues than any leadership memo.

Legal and privacy offices. Counsel’s job is to ask “What about the data?” Arrive with the answer. Bring the specifics of the tier you are buying, including the contractual commitment that inputs are not used for model training, the admin controls, the security authorizations such as FedRAMP or StateRAMP, and a first use case that touches no legally protected or case-level data at all. Scoping the first projects to non-sensitive work manages risk, and it also turns your legal review from a blocker into a partner who helps you sequence toward the harder use cases later.

IT and security. Security review is where most government AI efforts die, so it deserves the most preparation. IT and security teams are the most common reason agencies do not adopt AI, sometimes for good reasons and often for reflexive ones, and the reflexive no deserves a direct answer. The frontier tools in question are not consumer chatbots. On their enterprise tiers, Claude and ChatGPT carry contractual guarantees that your inputs are not used to train the model, and they have cleared federal security bars well beyond what a typical state agency requires. Claude is authorized at FedRAMP High through AWS GovCloud, ChatGPT runs in an authorized federal government cloud, and in 2025 the federal government offered both to federal agencies for one dollar.² The same class of tools is already at work inside NASA, the national laboratories, the National Institutes of Health, and the Treasury Department. If a tool is trusted there, the claim that it is too risky for your agency to use for summarizing public comments is hard to sustain. Start the security conversation from a presumption of yes, meaning “we are going to find a safe way to do this,” rather than a presumption of no. Put it on the governance agenda early, keep the first use case clear of legally protected and case-level data so the review is straightforward, and frame system integration as a separate, later decision with its own review. Your security team’s job is to help you do this safely. The decision about whether to do it at all belongs to agency leadership.

Funding and procurement. Small pilots usually fit within existing administrative budgets, and at a few thousand dollars a year they fall, in many states, under small-purchase thresholds that avoid competitive procurement. There is a federal dimension to keep in mind. If you later seek federal funding for larger AI investments in systems that support federally funded programs, advance approval documents and federal cost-allocation rules may apply. The specifics vary by program and

state, this paper does not offer guidance on them, and your budget and procurement staff should verify the right path before you commit. For the first pilots, keep them small enough that they do not trigger those processes at all.

Legislators and the public. The vendor pitches and legislative questions from the opening of this paper do not go away, and you need a posture for both. For legislators, the strongest position is specific and modest: “We are running small, low-risk internal pilots with human review, this is what the first one saved, and this is what we are trying next.” That answer works for the member who thinks you are moving too slowly and for the one who fears you are moving too fast. For vendors, the foundation above gives you a polite filter. Until your team has hands-on experience from a few pilots, you are not in a position to evaluate enterprise AI products well, which is a reason to start now and to wait on the big contracts.

What you can do this week

The only way to start is by starting, and starting is much smaller than it looks. A strategy, a budget line, and a dedicated team can all come later. To begin, you need an owner, a handful of licenses, one well-chosen use case, and the discipline to judge the result on its merits. The strategy, the governance, and the bigger ambitions all build better on three months of experience than on six months of planning.

A director could do all of the following within five business days:

- Name the owner, the one person accountable for running the learning loop.
- Ask your managers one question, “What recurring work eats the most time for the least judgment?”, and start the use-case list from their answers.
- Get the security conversation on the calendar with a slot on the next IT governance agenda.
- Price the pilot at five to ten business-tier licenses and confirm it against your small-purchase threshold.
- Pick a candidate first use case from your list and score it against the checklist.

Ninety days from now, you can have a working tool, a team that trusts it, a credible measure of what it saved, and an agency that is learning about AI faster than it is being pitched on it.

Notes & references

1. Christos Makridis, "AI Adoption Rapidly Growing in Public Sector," Gallup, March 11, 2026 (survey fielded Q4 2025, Oct. 30–Nov. 13, 2025). gallup.com/workplace/702983.
2. Claude foundation models reached FedRAMP High and DoD IL4/5 via Amazon Bedrock in AWS GovCloud in 2025. OpenAI's ChatGPT Enterprise is FedRAMP-authorized, with a self-hostable government build for higher impact levels. Under GSA "OneGov" agreements, Claude (Aug. 12, 2025) and ChatGPT Enterprise (Aug. 6, 2025) were each offered across federal agencies for \$1 for one year. Federal users named under "OpenAI for Government" include NASA, the national laboratories, the National Institutes of Health, and the Treasury Department. Sources: gsa.gov, anthropic.com, openai.com, and aws.amazon.com.

17A is a mission-driven consulting, talent, and technology firm for state and local government, working to put excellent people on the public problems that matter most. Learn more at group17a.com.

hello@group17a.com

77 Sands Street, Floor 6, Office 7067, Brooklyn, NY 11201